

# COSC 311: Introduction to Data Visualization and Interpretation

## (2) SQL

Dr. Shuangquan (Peter) Wang  
([spwang@salisbury.edu](mailto:spwang@salisbury.edu))

Department of Computer Science  
Salisbury University



# About this note

---

- The contents of this note refer to:
  - Prof. Jeff Ullman and Prof. Jennifer Widom's Teaching Materials (Stanford University)
    - <http://infolab.stanford.edu/~ullman/fcdb.html>
  - Prof. Fang Li's Teaching Materials (Shanghai Jiao Tong University)
    - <https://www.cs.sjtu.edu.cn/~li-fang/DB.htm>
  - Prof. Tim Finin's Teaching Materials (UMBC)
    - <https://www.csee.umbc.edu/courses/461/461.shtml>
  - Prof. Matei Zaharia's Teaching Materials (Stanford University)
    - <https://web.stanford.edu/class/cs245/>
  - SQL Tutorial
    - <https://www.w3schools.com/sql/default.asp>
  - Some pictures are from Google search

**Dissemination or sale of any part of this note is NOT permitted!**

# Content

---

1. Simple queries in SQL
2. Queries involving more than one relation
3. Subqueries
4. Aggregation

# Language SQL

---

- SQL is a standard database language, adopted by many commercial systems
  - How to query the database
  - How to make modifications on database
  - Transactions in SQL

# Why SQL?

---

- SQL is a very-high-level language.
  - Say “what to do” rather than “how to do it.”
  - Avoid a lot of data-manipulation details needed in procedural languages like C++ or Java.
- Database management system (DBMS) figures out the “best” way to execute query.
  - Called “query optimization”

# SQL: structured query language

---

- Components of language:
  - Schema definition, Data retrieval, Data modification, Indexes, Constraints, Views, Triggers, Transactions, authorization, etc.
- DDL = data definition language
- DML = data manipulation language
- Two forms of usage:
  - Interactive SQL (GUI, prompt)
  - Embedded SQL (Python, C, Java)

# Select-From-Where Statements

---

**SELECT** <desired attributes>

**FROM** <tuple variables or relation name>

**WHERE** <conditions>

**GROUP BY** <attributes>

Principal form

**HAVING** <conditions>

**ORDER BY** < list of attributes>

# Our Running Example

---

- All our SQL queries will be based on the following database schema.
  - Underline indicates key attributes.

Beers(name, manf)

Bars(name, addr, license)

Drinkers(name, addr, phone)

Likes(drinker, beer)

Sells(bar, beer, price)

Frequents(drinker, bar)



# Example: Query on one relation

---

- Using **Beers(name, manf)**, what beers are made by Anheuser-Busch?

```
SELECT name
```

```
FROM Beers
```

```
WHERE manf = 'Anheuser-Busch' ;
```

Note: single quotes for strings



name
Bud
Bud Lite
Michelob
...

The answer is a relation with a single attribute, name, and tuples with the name of each beer by Anheuser-Busch, such as Bud.

# Meaning of Single-Relation Query

---

- Begin with the relation in the **FROM** clause.
- Apply the selection indicated by the **WHERE** clause.
- Apply the extended projection indicated by the **SELECT** clause.

**SELECT** <desired attributes>

**FROM** <tuple variables or relation name>

**WHERE** <conditions>

# Operational Semantics

---

```
SELECT name
FROM Beers
WHERE manf = 'Anheuser-Busch';
```

name	manf
...	...
Bud	Anheuser-Busch
...	...

Tuple-variable  $t$   
loops over all  
tuples

Check if  
Anheuser-Busch

Include  $t.name$   
in the result, if so

# Operational Semantics --- General

---

- Think of a *tuple variable* visiting each tuple of the relation mentioned in FROM.
- Check if the “current” tuple satisfies the WHERE clause.
- If so, compute the attributes or expressions of the SELECT clause using the components of this tuple.

## \* In SELECT clauses

---

- When there is one relation in the FROM clause, \* in the SELECT clause stands for “all attributes of this relation.”
- **Example:** Using **Beers(name, manf):**

```
SELECT *  
FROM Beers  
WHERE manf = 'Anheuser-Busch';
```

# Result of Query:

---

name	manf
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
. . .	. . .

Now, the result has each of the attributes of Beers.

# Renaming Attributes

---

- If you want the result to have different attribute names, use “AS <new name>” to rename an attribute.

- **Example:** Using **Beers(name, manf):**

```
SELECT name AS beer, manf  
FROM Beers  
WHERE manf = 'Anheuser-Busch'
```

# Result of Query:

---

beer	manf
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
. . .	. . .



# Expressions in **SELECT** Clauses

---

- Any expression that makes sense can appear as an element of a **SELECT** clause.
- Example:** Using **Sells(bar, beer, price):**

```
SELECT bar, beer,  
        price*114 AS priceInYen  
FROM Sells;
```

bar	beer	priceInYen
Joe's	Bud	285
Sue's	Miller	342
...	...	...

# Complex Conditions in **WHERE** Clause

---

- Boolean operators AND, OR, NOT.
- Comparisons =, <>, <, >, <=, >=.
  - And many other operators that produce boolean-valued results.

# Example: Complex Condition

---

- Using `Sells(bar, beer, price)`, find the price Joe's Bar charges for Bud:

```
SELECT price
FROM Sells
WHERE bar = 'Joe's Bar' AND
       beer = 'Bud';
```

# Patterns

---

- WHERE clauses can have conditions in which a string is compared with a pattern, to see if it matches.
- Form:
  - <Attribute> LIKE <pattern> or
  - <Attribute> NOT LIKE <pattern>
- *Pattern* is a quoted string with % = “any string”; \_ = “any character.”

Note: SQL is case insensitive. Keywords like SELECT or AND can be written upper/lower case as you like. Only inside quoted strings does case matter.

# Example: LIKE

---

- Using `Drinkers(name, addr, phone)` find the drinkers with exchange 555:

```
SELECT name
FROM Drinkers
WHERE phone LIKE '%555- _ _ _ _';
```

## 2. Queries involving more than one relation

---

- Interesting queries often combine data from more than one relation.
- We can address several relations in one query by listing them all in the FROM clause.
- Distinguish attributes of the same name by “<relation>.<attribute>” .

# Example: Joining Two Relations

---

- Using relations **Likes(drinker, beer)** and **Frequents(drinker, bar)**, find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes, Frequents
WHERE bar = 'Joe's Bar' AND
      Frequents.drinker =
        Likes.drinker;
```

# Formal Semantics

---

- Almost the same as for single-relation queries:
  1. Start with the product of all the relations in the FROM clause.
  2. Apply the selection condition from the WHERE clause.
  3. Project onto the list of attributes and expressions in the SELECT clause.



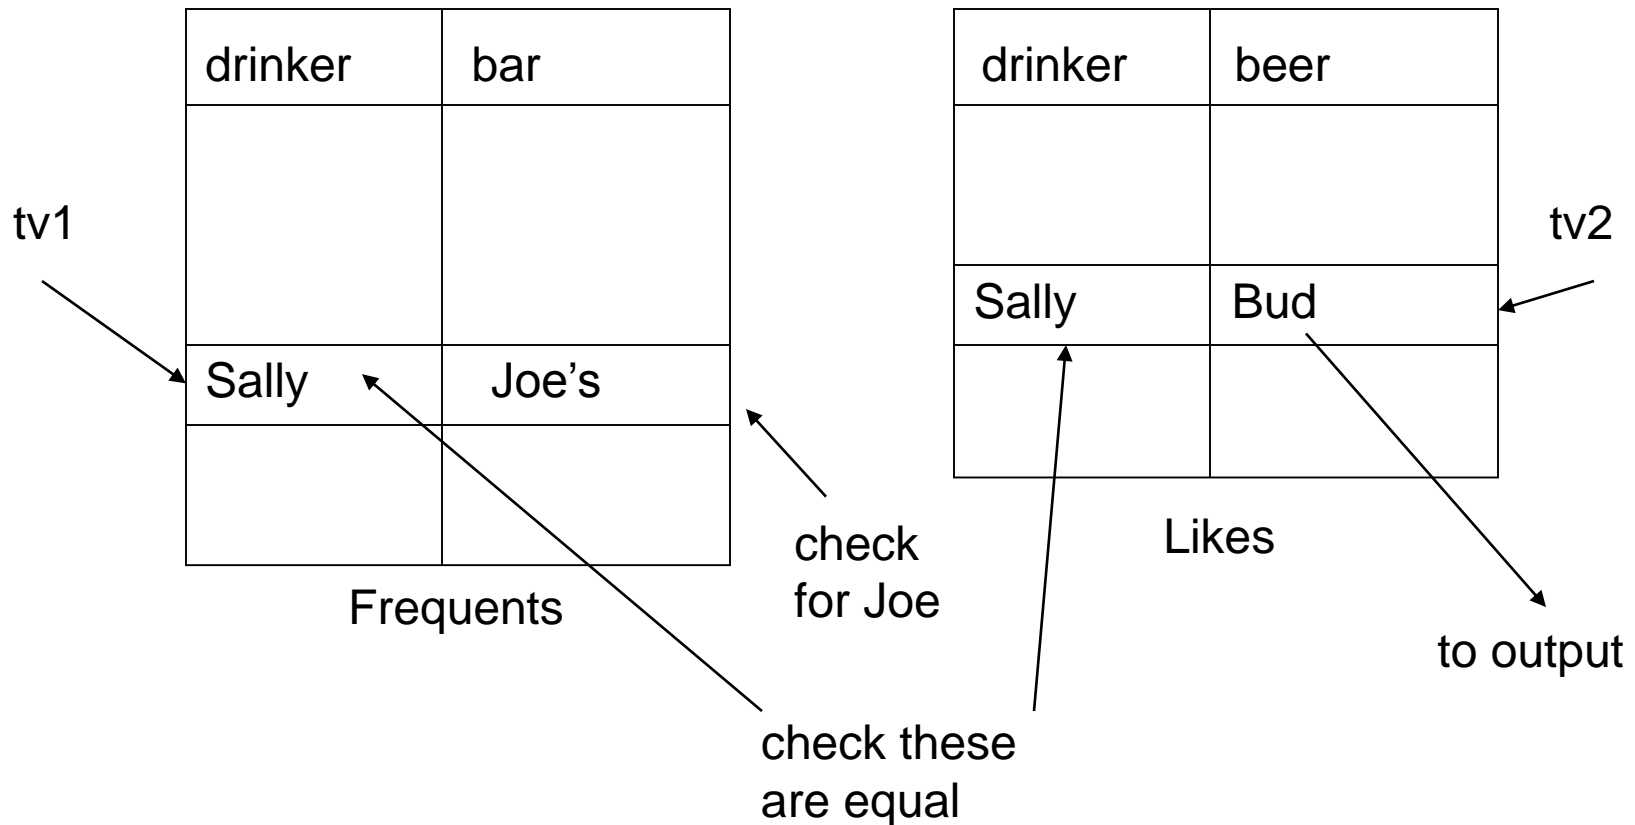
# Operational Semantics

---

- Imagine one tuple-variable for each relation in the FROM clause.
  - These tuple-variables visit each combination of tuples, one from each relation.
- If the tuple-variables are pointing to tuples that satisfy the WHERE clause, send these tuples to the SELECT clause.

# Example

```
SELECT beer
FROM Likes, Frequents
WHERE bar = 'Joe''s Bar' AND
      Frequents.drinker = Likes.drinker;
```



# 3. Subqueries

---

- A parenthesized SELECT-FROM-WHERE statement (*subquery*) can be used as a value in a number of places, including FROM and WHERE clauses.
- **Example:** in place of a relation in the FROM clause, we can use a subquery and then query its result.
  - Must use a tuple-variable to name tuples of the result.

# Example: Subquery in FROM

---

- Find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes, (SELECT drinker
FROM Frequent
WHERE bar = 'Joe's Bar') JD
WHERE Likes.drinker = JD.drinker;
```

Drinkers who  
frequent Joe's Bar

## Note:

- JD is a tuple-variable to name tuples of the result
- Parentheses around subquery are essential

# Subqueries That Return One Tuple

---

- If a subquery is guaranteed to produce one tuple, then the subquery can be used as a value.
  - Usually, the tuple has one component.
  - A run-time error occurs if there is no tuple or more than one tuple.

# Example: Single-Tuple Subquery

---


- Using **Sells(bar, beer, price)**, find the bars that serve Miller for the same price Joe charges for Bud.
- Two queries would surely work:
  1. Find the price Joe charges for Bud.
  2. Find the bars that serve Miller at that price.

# Query + Subquery Solution

---

```
SELECT bar
FROM Sells
WHERE beer = 'Miller' AND
      price = (
```

The price at  
which Joe  
sells Bud



```
SELECT price
FROM Sells
WHERE bar = 'Joe's Bar'
      AND beer = 'Bud');
```

Note the scoping rule: an attribute refers to the most closely nested relation with that attribute.

# The IN Operator

---

- `<tuple> IN (<subquery>)` is true if and only if the tuple is a member of the relation produced by the subquery.
  - Opposite: `<tuple> NOT IN (<subquery>)`.
- IN-expressions can appear in WHERE clauses.



# Example: IN

---

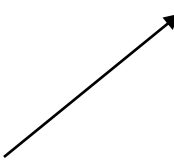
- Using **Beers(name, manf)** and **Likes(drinker, beer)**, find the name and manufacturer of each beer that Fred likes.

SELECT \*

FROM Beers

WHERE name IN (SELECT beer  
FROM Likes  
WHERE drinker = 'Fred');

The set of  
beers Fred  
likes



# Exercise: What is the difference?

---

```
SELECT a
FROM R, S
WHERE R.b = S.b;
```

a	b
1	2
3	4

R

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```

b	C
2	5
2	6

S

# Exercise: What is the difference?

```
SELECT a
FROM R, S
WHERE R.b = S.b;
```

Double loop,  
over the tuples  
of R and S

(1,2) with (2,5) and  
(1,2) with (2,6) both  
satisfy the condition; 1  
is output twice.

a	b
1	2
3	4

R

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```

One loop,  
over the  
tuples of R

Two 2's

(1,2) satisfies  
the condition; 1  
is output once.

b	C
2	5
2	6

S

# The Exists Operator

---

- EXISTS(<subquery>) is true if and only if the subquery result is not empty.
- **Example:** From **Beers(name, manf)** , find those beers that are the unique beer by their manufacturer.

# Example: EXISTS

```
SELECT name  
FROM Beers b1  
WHERE NOT EXISTS (
```

Notice scope rule: manf refers to closest nested FROM with a relation having that attribute.

Set of beers with the same manf as b1, but not the same beer

```
SELECT *  
FROM Beers  
WHERE manf = b1.manf AND  
      name <> b1.name);
```

Notice the SQL “not equals” operator

Note: A subquery that refers to values from a surrounding query is called a correlated subquery.

## 4. Aggregation

---

- SUM, AVG, COUNT, MIN, and MAX can be applied to a column in a SELECT clause to produce that aggregation on the column.
- Also, COUNT(\*) counts the number of tuples.

# Example: Aggregation

---

- From **Sells(bar, beer, price)**, find the average price of Bud:

```
SELECT AVG (price)
FROM Sells
WHERE beer = 'Bud' ;
```

# Eliminating Duplicates in an Aggregation

---

- Use DISTINCT inside an aggregation.
- **Example:** From `Sells(bar, beer, price)`, find the number of *different* prices charged for Bud:

```
SELECT COUNT (DISTINCT price)
FROM Sells
WHERE beer = 'Bud' ;
```

Note: DISTINCT may be used in any aggregation, but typically only makes sense with COUNT.



# Grouping

---

- We may follow a SELECT-FROM-WHERE expression by GROUP BY and a list of attributes.
- The relation that results from the SELECT-FROM-WHERE **is grouped** according to the values of all those attributes, and any aggregation is applied only **within each group**.

# Example: Grouping

---

- From **Sells(bar, beer, price)**, find the average price for each beer:

```
SELECT beer, AVG(price)
FROM Sells
GROUP BY beer;
```

beer	AVG(price)
Bud	2.33
...	...

# Example: Grouping

---

- From **Sells(bar, beer, price)** and **Frequents(drinker, bar)**, find for each drinker the average price of Bud at the bars they frequent:

```
SELECT drinker, AVG(price)
FROM Frequents, Sells
WHERE beer = 'Bud' AND
```

Compute all  
drinker-bar-  
price triples  
for Bud.

Frequents.bar = Sells.bar

```
GROUP BY drinker;
```

Then group  
them by  
drinker.

---

# Thanks

# Reading textbook

---

- A First Course in Database systems (Third Edition), by Jeff Ullman and Jennifer Widom.
  - Chapter 6